

# Linear Algebra for AI/ML

Jiaul Paik

# Latent Semantic Analysis of Text with SVD

- Term-document matrices are very large
- But the number of topics that people talk about is small (in some sense)
  - Clothes, movies, politics, ...
- Can we represent the term-document space by a lower dimensional latent space?

# Similarity $\rightarrow$ Clustering

- We can compute the similarity between two document vector representations  $\mathbf{x}_i$  and  $\mathbf{x}_j$  by  $\mathbf{x}_i \mathbf{x}_j^T$
- Let  $\mathbf{X} = [\mathbf{x}_1 \dots \mathbf{x}_N]$
- Then  $\mathbf{X}\mathbf{X}^T$  is a matrix of similarities
- $\mathbf{X}_{ij}$  is symmetric
- So  $\mathbf{X}\mathbf{X}^T = \mathbf{Q}\mathbf{\Lambda}\mathbf{Q}^T$
- So we can decompose this similarity space into a set of orthonormal basis vectors (given in  $\mathbf{Q}$ ) scaled by the eigenvalues in  $\mathbf{\Lambda}$

# Singular Value Decomposition

For an  $M \cdot N$  matrix  $\mathbf{A}$  of rank  $r$  there exists a factorization (Singular Value Decomposition = **SVD**) as follows:

$$\mathbf{A} = \mathbf{U} \mathbf{\Sigma} \mathbf{V}^T$$

The diagram illustrates the dimensions of the matrices in the SVD equation  $\mathbf{A} = \mathbf{U} \mathbf{\Sigma} \mathbf{V}^T$ . Below the equation, three boxes contain the dimensions:  $M \cdot M$ ,  $M \cdot N$ , and  $V \text{ is } N \cdot N$ . Arrows point from these boxes to the matrices  $\mathbf{U}$ ,  $\mathbf{\Sigma}$ , and  $\mathbf{V}^T$  respectively.

# Singular Value Decomposition

- Illustration of SVD dimensions and sparseness

$$\underbrace{\begin{bmatrix} * & * & * \\ * & * & * \\ * & * & * \\ * & * & * \\ * & * & * \end{bmatrix}}_A = \underbrace{\begin{bmatrix} * & * & * & * & * \\ * & * & * & * & * \\ * & * & * & * & * \\ * & * & * & * & * \\ * & * & * & * & * \end{bmatrix}}_U \underbrace{\begin{bmatrix} \bullet & & & & \\ & \bullet & & & \\ & & \bullet & & \\ & & & \bullet & \\ & & & & \bullet \end{bmatrix}}_{\Sigma} \underbrace{\begin{bmatrix} * & * & * \\ * & * & * \\ * & * & * \end{bmatrix}}_{V^T}$$

$$\underbrace{\begin{bmatrix} * & * & * & * & * \\ * & * & * & * & * \\ * & * & * & * & * \end{bmatrix}}_A = \underbrace{\begin{bmatrix} * & * & * \\ * & * & * \\ * & * & * \end{bmatrix}}_U \underbrace{\begin{bmatrix} \bullet & & & & \\ & \bullet & & & \\ & & \bullet & & \\ & & & \bullet & \\ & & & & \bullet \end{bmatrix}}_{\Sigma} \underbrace{\begin{bmatrix} * & * & * & * & * \\ * & * & * & * & * \\ * & * & * & * & * \\ * & * & * & * & * \\ * & * & * & * & * \end{bmatrix}}_{V^T}$$

# Low-rank Approximation

- Solution via SVD

*set smallest  $r-k$  singular values to zero*

$$\underbrace{\begin{bmatrix} * & * & * & * & * \\ * & * & * & * & * \\ * & * & * & * & * \end{bmatrix}}_{A_k} = \underbrace{\begin{bmatrix} * & * & * \\ * & * & * \\ * & * & * \end{bmatrix}}_U \underbrace{\begin{bmatrix} \bullet & & & & \\ & \bullet & & & \\ & & & & \end{bmatrix}}_{\Sigma} \underbrace{\begin{bmatrix} * & * & * & * & * \\ * & * & * & * & * \\ * & * & * & * & * \\ * & * & * & * & * \\ * & * & * & * & * \end{bmatrix}}_{V^T}$$

# Reduced SVD

- If we retain only  $k$  singular values, and set the rest to 0, then we don't need the matrix parts in color
- Then  $\Sigma$  is  $k \times k$ ,  $U$  is  $M \times k$ ,  $V^T$  is  $k \times N$ , and  $A_k$  is  $M \times N$
- This is referred to as the reduced SVD
- It is the convenient (space-saving) and usual form for computational applications

The diagram illustrates the Reduced SVD decomposition of a matrix  $A_k$  into three components:  $U$ ,  $\Sigma$ , and  $V^T$ .

On the left, the matrix  $A_k$  is shown as a 3x5 grid of stars. Below it is a brace labeled  $A_k$ .

An equals sign follows, leading to the product of three matrices:

- $U$ : A 3x3 grid of stars, with the third column highlighted in blue. Below it is a brace labeled  $U$ .
- $\Sigma$ : A 3x3 matrix with a blue L-shaped region containing two dots and a blue square at the bottom-right, and a yellow rectangular region to its right. Below it is a brace labeled  $\Sigma$ .
- $V^T$ : A 3x5 grid of stars, with the second row highlighted in blue and the bottom two rows highlighted in yellow. Below it is a brace labeled  $V^T$ .

# Latent Semantic Indexing via the SVD



# What it is

- From term-doc matrix  $A$ , we compute the approximation  $A_k$ .
- There is a row for each term and a column for each doc in  $A_k$
- Thus docs live in a space of  $k \ll r$  dimensions
  - These dimensions are not the original axes
- But why?

# Text Mining/NLP: Vector Space Model

- **Automatic** selection of index terms
- **Partial matching** of queries and documents (dealing with the case where no document contains all search terms)
- **Ranking** according to **similarity score** (dealing with large result sets)
- **Term weighting** schemes (improves retrieval performance)
- **Various extensions**
  - Document clustering
  - Relevance feedback (modifying query vector)
- **Geometric foundation**

# Problems with Lexical Semantics

- **Ambiguity and association in natural language**
  - **Polysemy**: Words often have a **multitude of meanings** and different types of usage (more severe in very heterogeneous collections).
  - The vector space model is unable to discriminate between different meanings of the same word.

$$\text{sim}_{\text{true}}(d, q) < \cos(\angle(\vec{d}, \vec{q}))$$

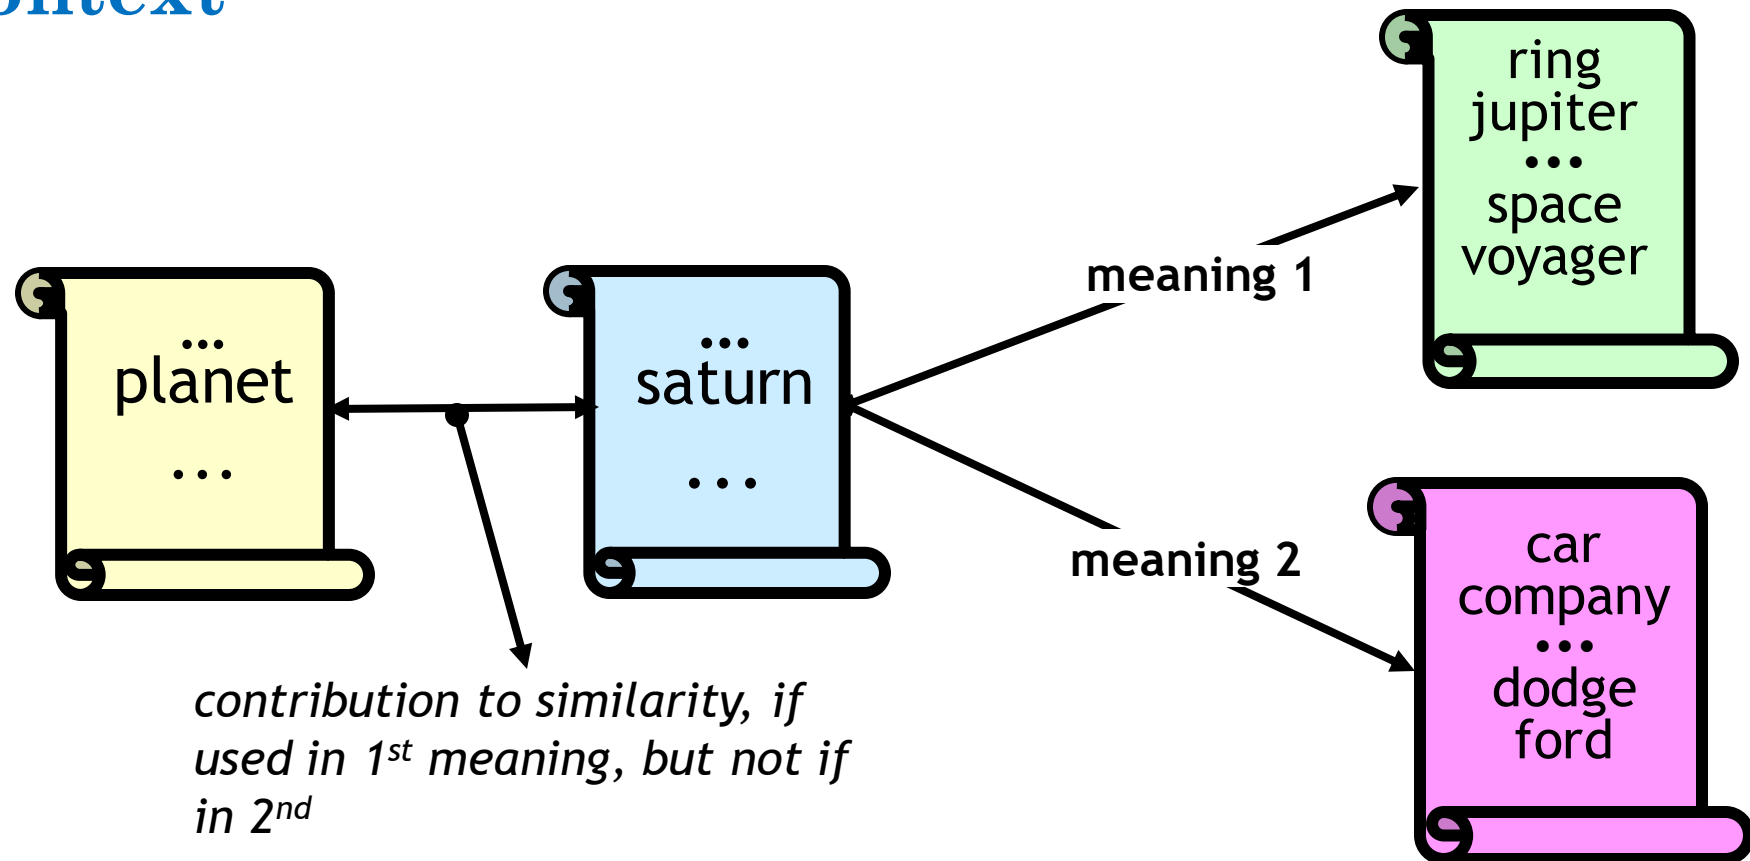
# Problems with Lexical Semantics

- **Synonymy**: Different terms may have an identical or a similar meaning (weaker: words indicating the same topic).
- No associations between words are made in the vector space representation.

$$\text{sim}_{\text{true}}(d, q) > \cos(\angle(\vec{d}, \vec{q}))$$

# Polysemy and Context

- Document similarity on single word level: **polysemy and context**

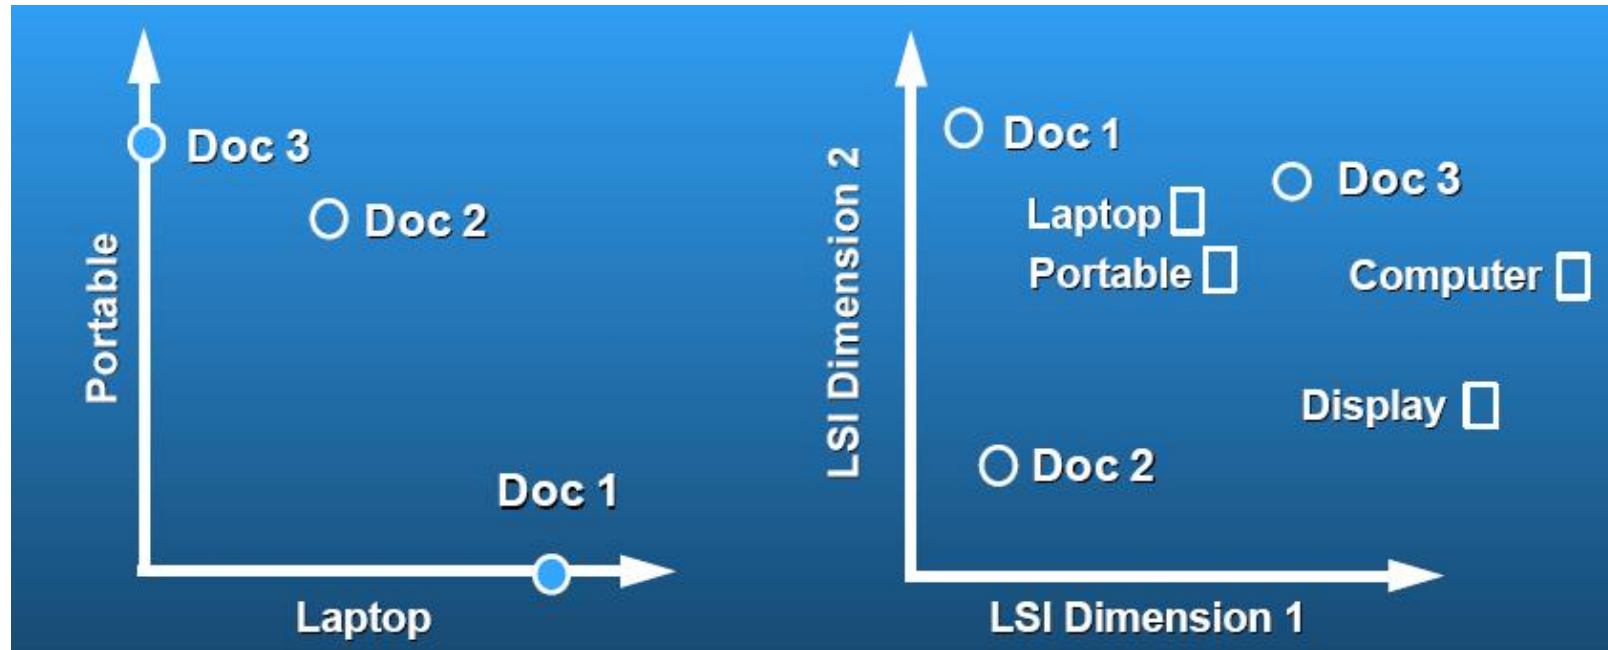


# Latent Semantic Analysis (LSA)

- **Perform a low-rank approximation of document-term matrix**
- **General idea**
  - Map documents (and terms) to a **low-dimensional** representation.
  - Design a mapping such that the low-dimensional space reflects **semantic associations** (latent semantic space).
  - Compute document similarity based on the **inner product** in this **latent semantic space**

# Latent Semantic Analysis

- **Latent semantic space:** illustrating example



*courtesy of Susan Dumais*

# LSA Example



# LSA Example

- A simple example term-document matrix (binary)

$C$	$d_1$	$d_2$	$d_3$	$d_4$	$d_5$	$d_6$
ship	1	0	1	0	0	0
boat	0	1	0	0	0	0
ocean	1	1	0	0	0	0
wood	1	0	0	1	1	0
tree	0	0	0	1	0	1

# LSA Example

- Example of  $C = U\Sigma V^T$ : The matrix  $U$

$U$	1	2	3	4	5
ship	−0.44	−0.30	0.57	0.58	0.25
boat	−0.13	−0.33	−0.59	0.00	0.73
ocean	−0.48	−0.51	−0.37	0.00	−0.61
wood	−0.70	0.35	0.15	−0.58	0.16
tree	−0.26	0.65	−0.41	0.58	−0.09

# LSA Example

- Example of  $C = U\Sigma V^T$ : The matrix  $\Sigma$

$\Sigma$	1	2	3	4	5
1	2.16	0.00	0.00	0.00	0.00
2	0.00	1.59	0.00	0.00	0.00
3	0.00	0.00	1.28	0.00	0.00
4	0.00	0.00	0.00	1.00	0.00
5	0.00	0.00	0.00	0.00	0.39

# LSA Example

- Example of  $C = U\Sigma V^T$ : The matrix  $V^T$

$V^T$	$d_1$	$d_2$	$d_3$	$d_4$	$d_5$	$d_6$
1	-0.75	-0.28	-0.20	-0.45	-0.33	-0.12
2	-0.29	-0.53	-0.19	0.63	0.22	0.41
3	0.28	-0.75	0.45	-0.20	0.12	-0.33
4	0.00	0.00	0.58	0.00	-0.58	0.58
5	-0.53	0.29	0.63	0.19	0.41	-0.22

# LSA Example: Reducing the dimension

$U$	1	2	3	4	5	
ship	−0.44	−0.30	0.00	0.00	0.00	
boat	−0.13	−0.33	0.00	0.00	0.00	
ocean	−0.48	−0.51	0.00	0.00	0.00	
wood	−0.70	0.35	0.00	0.00	0.00	
tree	−0.26	0.65	0.00	0.00	0.00	
$\Sigma_2$	1	2	3	4	5	
1	2.16	0.00	0.00	0.00	0.00	
2	0.00	1.59	0.00	0.00	0.00	
3	0.00	0.00	0.00	0.00	0.00	
4	0.00	0.00	0.00	0.00	0.00	
5	0.00	0.00	0.00	0.00	0.00	
$V^T$	$d_1$	$d_2$	$d_3$	$d_4$	$d_5$	$d_6$
1	−0.75	−0.28	−0.20	−0.45	−0.33	−0.12
2	−0.29	−0.53	−0.19	0.63	0.22	0.41
3	0.00	0.00	0.00	0.00	0.00	0.00
4	0.00	0.00	0.00	0.00	0.00	0.00
5	0.00	0.00	0.00	0.00	0.00	0.00

# Original matrix $C$ vs. $C_2 = U\Sigma_2V^T$


$C$	$d_1$	$d_2$	$d_3$	$d_4$	$d_5$	$d_6$
ship	1	0	1	0	0	0
boat	0	1	0	0	0	0
ocean	1	1	0	0	0	0
wood	1	0	0	1	1	0
tree	0	0	0	1	0	1

$C_2$	$d_1$	$d_2$	$d_3$	$d_4$	$d_5$	$d_6$
ship	0.85	0.52	0.28	0.13	0.21	-0.08
boat	0.36	0.36	0.16	-0.20	-0.02	-0.18
ocean	1.01	0.72	0.36	-0.04	0.16	-0.21
wood	0.97	0.12	0.20	1.03	0.62	0.41
tree	0.12	-0.39	-0.08	0.90	0.41	0.49

# Why the reduced dimension matrix is better

- **Similarity of d2 and d3 in the original space: 0.**
- **Similarity of d2 and d3 in the reduced space:**
  - $0.52 * 0.28 + 0.36 * 0.16 + 0.72 * 0.36 + 0.12 * 0.20 + -0.39 * -0.08 \approx 0.52$

# Simplistic picture





# LSI has many other applications

- **In many settings in pattern recognition and retrieval, we have a feature-object matrix.**
  - For text, the terms are features and the docs are objects.
  - Could be opinions and users ...
  - This matrix may be redundant in dimensionality.
  - Can work with low-rank approximation.
  - If entries are missing (e.g., users' opinions), can recover if dimensionality is low.
- **Powerful general analytical technique**
  - Close, principled analog to clustering methods.

Thank you!