

# Tutorial 10: Heaps

Algorithms - I

By Kanishk Singh

# Table of Contents

- ① Previous Year's Lab Assignment
  - Problem
  - Hints and Solution
  - Implementations

# Problem Statement

- Consider a refrigerator servicing facility that can handle just one refrigerator at a time. Suppose that on a particular day there are  $n$ -refrigerators awaiting repair. Arrival of refrigerator- $i$  is specified by a start time  $s_i$ , and a processing time  $p_i$ . You are allowed to suspend the repairing of refrigerator and resume it later. We will refer to each refrigerator servicing as a job.

# Problem Statement

- Consider a refrigerator servicing facility that can handle just one refrigerator at a time. Suppose that on a particular day there are  $n$ -refrigerators awaiting repair. Arrival of refrigerator- $i$  is specified by a start time  $s_i$ , and a processing time  $p_i$ . You are allowed to suspend the repairing of refrigerator and resume it later. We will refer to each refrigerator servicing as a job.
- A schedule specifies for each unit time interval, the unique job that is run during that time interval. In a feasible schedule, every job  $J_i$  has to be run for exactly  $p_i$  time units (*jobLength*) after time  $s_i$  (*startTime*).

# Problem Statement

- Consider a refrigerator servicing facility that can handle just one refrigerator at a time. Suppose that on a particular day there are  $n$ -refrigerators awaiting repair. Arrival of refrigerator- $i$  is specified by a start time  $s_i$ , and a processing time  $p_i$ . You are allowed to suspend the repairing of refrigerator and resume it later. We will refer to each refrigerator servicing as a job.
- A schedule specifies for each unit time interval, the unique job that is run during that time interval. In a feasible schedule, every job  $J_i$  has to be run for exactly  $p_i$  time units (*jobLength*) after time  $s_i$  (*startTime*).
- The completion time  $C_i$  for job  $J_i$  is the earliest time when  $J_i$  has been run for  $p_i$  time units. You are asked to implement a scheduler that minimizes :

$$\sum_{j=1}^n C_j$$

# Understand the plot

We want to schedule in an online fashion and it would look as follows :

$t = 0$

while there are jobs left not completely scheduled

- Among those jobs  $J_i$  such that  $s_i \leq t$ .
- pick a job  $J_m$  to schedule at time  $t$  according to some rule.
- increment  $t$ .

# A Working Example

- You are given the number of jobs and *jobId*, *startTime* and *jobLength* values for each job.
- Number of jobs :  $n = 4$
- The input format is given below :

jobId	startTime	jobLength
1	0	21
2	1	3
3	2	6
4	3	2

# A Working Example

- You are given the number of jobs and *jobId*, *startTime* and *jobLength* values for each job.
- Number of jobs :  $n = 4$
- The input format is given below :

jobId	startTime	jobLength
-------	-----------	-----------

1	0	21
---	---	----

2	1	3
---	---	---

3	2	6
---	---	---

4	3	2
---	---	---

- Jobs scheduled at each timestep are:
- 1 2 2 2 4 4 3 3 3 3 3 1



- You are also asked to print the average turnaround time.

- You are also asked to print the average turnaround time.
- Let's denote the *startTime* value for a job by  $s_i$  and the time when it first starts running on the processor by  $k_i$ . The average turnaround time is given by :

$$\frac{\sum_{i=1}^n (k_i - s_i)}{n}$$

- You are also asked to print the average turnaround time.
- Let's denote the *startTime* value for a job by  $s_i$  and the time when it first starts running on the processor by  $k_i$ . The average turnaround time is given by :

$$\frac{\sum_{i=1}^n (k_i - s_i)}{n}$$

- The individual turnaround times for the jobs are 0, 0, 4 and 1 respectively.

Hence Average Turnoaround Time is : 1.25

# Table of Contents

## ① Previous Year's Lab Assignment

Problem

Hints and Solution

Implementations

# Can you think of an approach?

Have 2 – 3 minutes to think.

# Greedy Approach

- Let  $y_{i,t}$  be the total time that job  $J_i$  has been run before time  $t$ .

# Greedy Approach

- Let  $y_{i,t}$  be the total time that job  $J_i$  has been run before time  $t$ .
- Pick  $J_m$  to be a job that has minimal remaining processing time, that is, that has minimal  $p_i - y_{i,t}$  (remLength).

# Greedy Approach

- Let  $y_{i,t}$  be the total time that job  $J_i$  has been run before time  $t$ .
- Pick  $J_m$  to be a job that has minimal remaining processing time, that is, that has minimal  $p_i - y_{i,t}$  (remLength).
- Ties may be broken based on unique jobid, so that lower jobid is preferred.



# Greedy Approach

- An efficient implementation of this greedy algorithm can be achieved by the scheduler maintaining a priority queue of currently remaining jobs.

# Greedy Approach

- An efficient implementation of this greedy algorithm can be achieved by the scheduler maintaining a priority queue of currently remaining jobs.
- When a job is executed on the processor, its remLength value is decremented by 1 after each timestep.

# Greedy Approach

- An efficient implementation of this greedy algorithm can be achieved by the scheduler maintaining a priority queue of currently remaining jobs.
- When a job is executed on the processor, its `remLength` value is decremented by 1 after each timestep.
- When the currently executing job finishes (`remLength` value 0), the scheduler removes the job with the least `remLength` value from the queue and schedules it to run on the processor.

# Greedy Approach

- An efficient implementation of this greedy algorithm can be achieved by the scheduler maintaining a priority queue of currently remaining jobs.
- When a job is executed on the processor, its remLength value is decremented by 1 after each timestep.
- When the currently executing job finishes (remLength value 0), the scheduler removes the job with the least remLength value from the queue and schedules it to run on the processor.
  - This job now becomes the currently executing job.

# Greedy Approach

- When a new job  $x$  arrives, its job duration is checked with the `remLength` value of the currently executing job  $y$ . Two cases are possible:

# Greedy Approach

- When a new job  $x$  arrives, its job duration is checked with the `remLength` value of the currently executing job  $y$ . Two cases are possible:
  - If the job duration of  $x$  is less than the `remLength` value of  $y$ ,  $y$  (with its current `remLength` value) is added to the scheduling queue and  $x$  becomes the currently executing job.

# Greedy Approach

- When a new job  $x$  arrives, its job duration is checked with the `remLength` value of the currently executing job  $y$ . Two cases are possible:
  - If the job duration of  $x$  is less than the `remLength` value of  $y$ ,  $y$  (with its current `remLength` value) is added to the scheduling queue and  $x$  becomes the currently executing job.
  - Otherwise,  $x$  is added to the scheduling queue.

# Greedy Approach

- When a new job  $x$  arrives, its job duration is checked with the `remLength` value of the currently executing job  $y$ . Two cases are possible:
  - If the job duration of  $x$  is less than the `remLength` value of  $y$ ,  $y$  (with its current `remLength` value) is added to the scheduling queue and  $x$  becomes the currently executing job.
  - Otherwise,  $x$  is added to the scheduling queue.
- If there are two jobs with the same `remLength` value, the one with the lower job id is to be chosen.



# Let's Rework The Example

jobId	startTime	jobLength
1	0	21
2	1	3
3	2	6
4	3	2

# Continued...

# Table of Contents

## ① Previous Year's Lab Assignment

Problem

Hints and Solution

**Implementations**

# Scheduling queue Implementation

- First, let's define a struct to represent a job.

```
#define MAXSIZE
typedef struct job{
    int jobId;
    int startTime;
    int jobLength;
    int remLength;
} job;
```

# Scheduling queue Implementation

- First, let's define a struct to represent a job.

```
#define MAXSIZE
typedef struct job{
    int jobId;
    int startTime;
    int jobLength;
    int remLength;
} job;
```

- heap of jobs then could be represented as:

```
typedef struct heap{
    job list[MAX_SIZE];
    int numJobs;
} heap;
```

# Scheduling queue Implementation

- First, let's define a struct to represent a job.

```
#define MAXSIZE
typedef struct job{
    int jobId;
    int startTime;
    int jobLength;
    int remLength;
} job;
```

- heap of jobs then could be represented as:

```
typedef struct heap{
    job list[MAX_SIZE];
    int numJobs;
} heap;
```

- Following are the helper functions to implement the queue.

```
void initHeap(heap *H): (sets numJobs to 0)
void insertJob(heap *H, job j): (inserts the job j in heap)
int extractMinJob(heap *H, job *j): (returns minimum element)
void decreaseKey(heap *H, job *j): (decreases key value of job j)
```

# Scheduling queue Implementation

- First, let's define a struct to represent a job.

```
#define MAXSIZE
typedef struct job{
    int jobId;
    int startTime;
    int jobLength;
    int remLength;
} job;
```

- heap of jobs then could be represented as:

```
typedef struct heap{
    job list[MAX_SIZE];
    int numJobs;
} heap;
```

- Following are the helper functions to implement the queue.

```
void initHeap(heap *H): (sets numJobs to 0)
void insertJob(heap *H, job j): (inserts the job j in heap)
int extractMinJob(heap *H, job *j): (returns minimum element)
void decreaseKey(heap *H, job *j): (decreases key value of job j)
```

---

## Algorithm 1 ScheduleJobs

---

```
1: initHeap( $H$ )
2: sort(jobList)
3:  $time \leftarrow 0$ 
4: while jobLeft do
5:   if  $time == someJob.startTime$  then
6:     insertJob( $H$ , someJob)
7:   end if
8:   currJob  $\leftarrow$  extractMinJob( $H$ , temp)
9:    $currJob.remLength \leftarrow currJob.remLength - 1$ 
10:   $time \leftarrow time + 1$ 
11: end while
```

---



# What if there is a job dependency

- Now suppose that all the jobs are not independent, and there exist some pairs of jobs  $(x, y)$  such that if job  $x$  finishes before  $y$  starts, it can reduce the running time (remLength value) of  $y$  by 50%, if it has not started yet.

# What if there is a job dependency

- Now suppose that all the jobs are not independent, and there exist some pairs of jobs  $(x, y)$  such that if job  $x$  finishes before  $y$  starts, it can reduce the running time (remLength value) of  $y$  by 50%, if it has not started yet.
- Relationship is one way only,  $y$  depends on  $x$  does not imply  $x$  depends on  $y$ , so  $(x, y)$  and  $(y, x)$  are different pairs.

# What if there is a job dependency

- Now suppose that all the jobs are not independent, and there exist some pairs of jobs  $(x, y)$  such that if job  $x$  finishes before  $y$  starts, it can reduce the running time (remLength value) of  $y$  by 50%, if it has not started yet.
- Relationship is one way only,  $y$  depends on  $x$  does not imply  $x$  depends on  $y$ , so  $(x, y)$  and  $(y, x)$  are different pairs.
- Your job is to schedule the jobs using the same policy, subject to the change mentioned above. However, note that in this case, the duration of a job already in the heap can change in the middle if another job that it depends on finishes.

# Changes Required?

- We need to decrease the key but how do we know the **position in the heap** ?

# Changes Required?

- We need to decrease the key but how do we know the **position in the heap** ?
- We can do that by storing the position of particular jobID in an array which would be updated after calling the heap functions.

```
typedef struct heap{  
    job list[MAX_SIZE];  
    int jobPos[MAX_SIZE];  
    int numJobs;  
} heap;
```

# Changes Required?

- We need to decrease the key but how do we know the **position in the heap** ?
- We can do that by storing the position of particular jobID in an array which would be updated after calling the heap functions.

```
typedef struct heap{  
    job list[MAX_SIZE];  
    int jobPos[MAX_SIZE];  
    int numJobs;  
} heap;
```

- Implement a helper function *decreaseKey* that would change the *remLength* value to *remLength/2*.

```
void decreaseKey(newheap *H, int jid);
```

# Changes Required?

- We need to decrease the key but how do we know the **position in the heap** ?
- We can do that by storing the position of particular jobID in an array which would be updated after calling the heap functions.

```
typedef struct heap{  
    job list[MAX_SIZE];  
    int jobPos[MAX_SIZE];  
    int numJobs;  
} heap;
```

- Implement a helper function *decreaseKey* that would change the *remLength* value to *remLength/2*.

```
void decreaseKey(newheap *H, int jid);
```

- Let's see the modified scheduler!!

---

## Algorithm 2 ScheduleJobs

---

```
1:  $\text{initHeap}(H)$ 
2:  $\text{sort}(\text{jobList})$ 
3: while  $\text{jobLeft}$  do
4:   if  $\text{time} == \text{someJob.startTime}$  then
5:      $\text{insertJob}(H, \text{someJob})$ 
6:   end if
7:    $\text{currJob} \leftarrow \text{extractMinJob}(H, \text{temp})$ 
8:    $\text{currJob.remLength} \leftarrow \text{currJob.remLength} - 1$ 
9:   if  $\text{currJob.remLength} \leq 0$  then
10:     $\text{someJob} \leftarrow \text{dependsOn}(\text{currJob})$ 
11:     $\text{decreaseKey}(H, \text{someJob})$ 
12:   end if
13:    $\text{time} \leftarrow \text{time} + 1$ 
14: end while}
```



# More Problems

Solve these problems on your own -

- [Monk and Multiplication](#)
- [Killing Zombies](#)
- [Monk And Champions League](#)

# Thank You!

Any Further Questions?