

Lecture 04

How to Implement a Dictionary?

- Sequences
 - ordered
 - unordered
- Binary Search Trees
- **Hash tables**

Hashing

- Another important and widely useful technique for implementing dictionaries
- Constant time per operation (on the average)
- Worst case time proportional to the size of the set for each operation (just like array and chain implementation)

Hashing - Basic Idea

- Use *hash function* to map keys into positions in a *hash table*

Ideally

- If element e has key k and h is hash function, then e is stored in position $h(k)$ of table
- To search for e , compute $h(k)$ to locate position. If no element, dictionary does not contain e .

Hash function example

- elements = Integers
- $h(i) = i \% 10$
- insert 41, 34, 7, and 18
- constant-time lookup:
 - just look at $i \% 10$ again later
- Hash tables have no ordering information!
 - Expensive to do following:
 - getMin, getMax, removeMin, removeMax,
 - the various ordered traversals
 - printing items in sorted order

0	
1	41
2	
3	
4	34
5	
6	
7	7
8	18
9	

Hashing Operations

- **Search**
 - looks for key k
- **Insert**
 - first searches for a slot, then inserts
- **Delete**
 - Cannot just turn the slot containing the key we want to delete to contain NIL. Why?

Hashing Analysis

- Analysis

- $O(b)$ time to initialize hash table (b number of positions or buckets in hash table)
- $O(1)$ time to perform *insert*, *remove*, *search*

- Reality

- Works for implementing dictionaries, but many applications have key ranges that are too large to have 1-1 mapping between buckets and keys!
- Example:
 - Suppose key can take on values from 0 .. 65,535 (2 byte unsigned int)
 - Expect $\approx 1,000$ records at any given time
 - Impractical to use hash table with 65,536 slots!

Hash Collisions

- **Collision:** the event that two hash table elements map into the same slot in the array
 - example: insert 41, 34, 7, 18, then 21
 - 21 hashes into the same slot as 41!
- **Resolution:**
 - How can we choose the hash function to minimize collisions?
 - What do we do about collisions when they occur?

0	
1	21
2	
3	
4	34
5	
6	
7	7
8	18
9	

Collision Resolution Policies

- Two classes:
 1. Closed hashing / open addressing
 2. Open hashing / separate chaining
- Difference has to do with whether collisions are stored *outside the table* (open hashing) or whether collisions result in storing one of the records at *another slot in the table* (closed hashing)

Open Addressing


- **Concept:**

- Store all n keys in the m slots of the hash table itself.
- Each slot contains either a key or NIL.
- To **search** for key k :
 - Examine slot $h(k)$. Examining a slot is known as a **probe**.
 - If slot $h(k)$ contains key k , the search is successful. If the slot contains NIL, the search is unsuccessful.
 - There's a third possibility: **slot $h(k)$ contains a key that is not k .**
 - Compute the index of some other slot, based on k and which probe we are on.
 - Keep probing until we either find key k or we find a slot holding NIL.

Advantages: Avoids pointers; so less code, and we can dedicate the memory to the table.

What can you say about the load factor $\alpha = n/m$?

Open addressing - issue



Closed Hashing

- Associated with closed hashing is a *rehash strategy*:
“If we try to place x in bucket $h(x)$ and find it occupied, find alternative location $h_1(x)$, $h_2(x)$, etc. Try each in order, if none empty table is full,”
- $h(x)$ is called *home bucket*
- Simplest rehash strategy is called *linear hashing*
$$h_i(x) = (h(x) + i) \% D$$
- In general, our collision resolution strategy is to generate a sequence of hash table slots (probe sequence) that can hold the record; test each slot until find empty one (probing)

Importance of Good Hash Functions

- Recall the assumption of *simple uniform hashing*:
 - Any key is equally likely to hash into any of the slots, independent of where any other key hashes to.
 - $O(1)$ time to compute $h(k)$.
- Hash values should be independent of any patterns that might exist in the data.
 - E.g. If each key is drawn independently from U according to a probability distribution P , we want
$$\text{for all } j \in [0 \dots m-1], \sum_{k:h(k)=j} P(k) = 1/m$$
- Often use heuristics, based on the domain of the keys, to create a hash function that performs well.

Two examples only

- **Division method**

- Map each key k into one of the m slots by taking the remainder of k divided by m .
$$h(k) = k \bmod m$$
- Example: $m = 31$ and $k = 78 \Rightarrow h(k) = 16$.
- Advantage: Fast, since requires just one division operation.
- Disadvantage: For some values, such as $m=2^p$, the hash depends on just a subset of the bits of the key.
- Note: Primes are good, if not too close to power of 2 (or 10).

- **Multiplication method**

- Map each key k to one of the m slots indicated by the fractional part of k times a chosen real $0 < A < 1$.

$$h(k) = \lfloor m (kA \bmod 1) \rfloor = \lfloor m (kA - \lfloor kA \rfloor) \rfloor$$

- Example: $m = 1000$, $k = 123$, $A \approx 0.6180339887...$

$$\begin{aligned} h(k) &= \lfloor 1000(123 \cdot 0.6180339887 \bmod 1) \rfloor \\ &= \lfloor 1000 \cdot 0.0181... \rfloor = 18. \end{aligned}$$

- Disadvantage: A bit slower than the division method.
- Advantage: Value of m is not critical.

Homework

**Implement these
two techniques.**

Example Linear (Closed) Hashing

- $D=8$, keys a, b, c, d have hash values $h(a)=3$, $h(b)=0$, $h(c)=4$, $h(d)=3$

⊕ Where do we insert d ? 3 already filled

⊕ Probe sequence using linear hashing:

$$h_1(d) = (h(d)+1)\%8 = 4\%8 = 4$$

$$h_2(d) = (h(d)+2)\%8 = 5\%8 = 5^*$$

$$h_3(d) = (h(d)+3)\%8 = 6\%8 = 6$$

etc.

7, 0, 1, 2

⊕ Wraps around the beginning of the table!

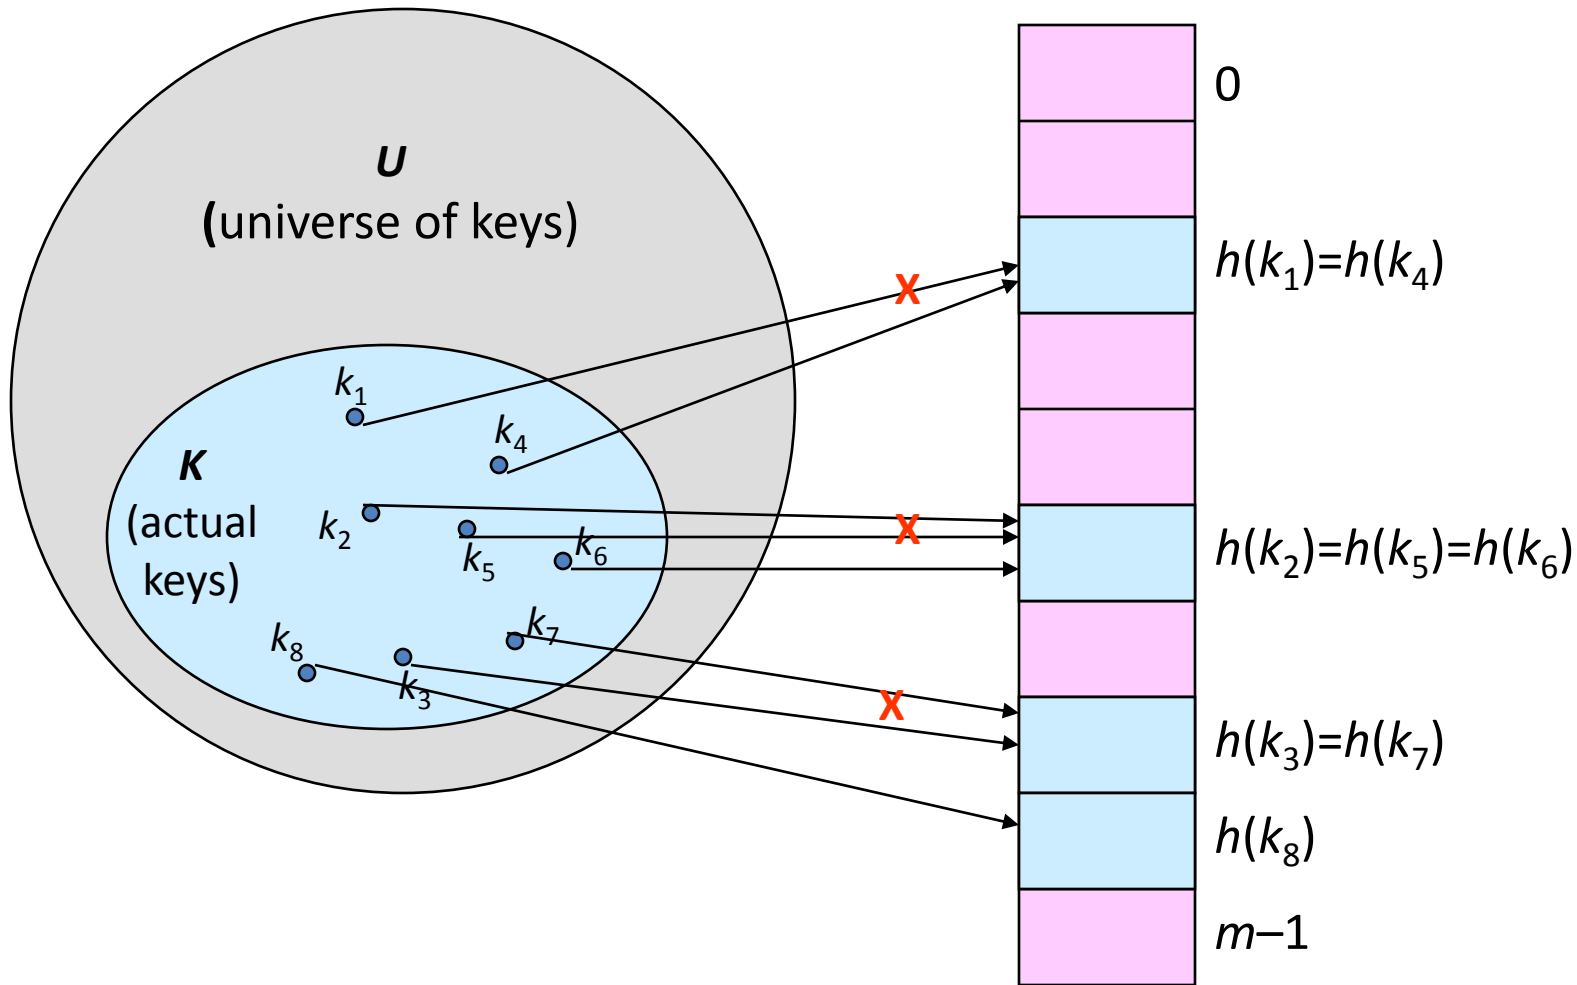
0	b
1	
2	
3	a
4	c
5	d
6	
7	

Performance Analysis


Computational complexity for initialization?

Computational complexity for insertion / search?

Collision Resolution by Chaining



Collision Resolution by Chaining



Hashing with Chaining

Dictionary Operations:

- Chained-Hash-Insert (T, x)
 - Insert x at the head of list $T[h(\text{key}[x])]$.
 - Worst-case complexity: $O(1)$.
- Chained-Hash-Search (T, k)
 - Search an element with key k in list $T[h(k)]$.
 - Worst-case complexity: proportional to length of list.
- Chained-Hash-Delete (T, x)
 - Delete x from the list $T[h(\text{key}[x])]$.
 - Worst-case complexity: search time + $O(1)$.
 - Need pointer to preceding element, or a doubly-linked list.

Analysis of Chained-Hash-Search

- ✓ **Worst-case search time:** time to compute $h(k) + \Theta(n)$.
- ✓ **Average time:** depends on how h distributes keys among slots.
 - ✓ **Assumptions:**
 - *Simple uniform hashing:* Any key is equally likely to hash into any of the slots, independent of where any other key hashes to.
 - $O(1)$ time to compute $h(k)$.
 - ✓ **Define** Load factor $\alpha = n/m$ = average # of keys per slot.
 - n – number of keys stored in the hash table.
 - m – number of slots = # linked lists.

Implications for separate chaining

- If $n = O(m)$, then load factor $\alpha = n/m = O(m)/m = O(1)$.
- Deletion takes $O(1)$ worst-case time if you have a pointer to the preceding element in the list.
- Hence, for hash tables with chaining, all dictionary operations take $O(1)$ time on average, given the assumptions of simple uniform hashing and $O(1)$ time hash function evaluation.
- Extra memory needed for linked list pointers.
- Can we satisfy the simple uniform hashing assumption?

Probe Sequence

- Sequence of slots examined during a key search constitutes a *probe sequence*.
- Probe sequence must be a permutation of the slot numbers.
 - We examine every slot in the table, if we have to.
 - We don't examine any slot more than once.
- One way to think of it: extend hash function to:
 - $h : U \times \underbrace{\{0, 1, \dots, m - 1\}}_{\text{probe number}} \rightarrow \underbrace{\{0, 1, \dots, m - 1\}}_{\text{slot number}}$

Universe of Keys

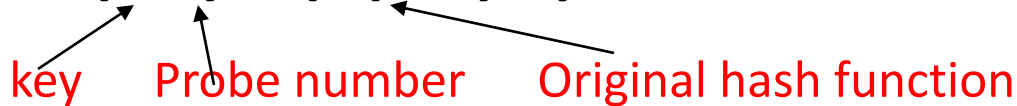
Computing Probe Sequences

- The ideal situation is *uniform hashing*:
 - Generalization of simple uniform hashing.
 - Each key is equally likely to have any of the $m!$ permutations of $\langle 0, 1, \dots, m-1 \rangle$ as its probe sequence.
 - It is hard to implement true uniform hashing.
- Approximate with techniques that guarantee to probe a permutation of $[0 \dots m-1]$, even if they don't produce all $m!$ probe sequences
 - Linear Probing.
 - Quadratic Probing.
 - Double Hashing.

Linear Probing

- $h(k, i) = (h(k, 0) + i) \bmod m$

key Probe number Original hash function



- The initial probe determines the entire probe sequence.
- Suffers from *primary clustering*:
 - Long runs of occupied sequences build up.
 - Long runs tend to get longer, since an empty slot preceded by i full slots gets filled next with probability $(i+1)/m$.

Clustering problem

- **Clustering:** nodes being placed close together by probing, which degrades hash table's performance
 - add 89, 18, 49, 58, 9
 - now searching for the value 28 will have to check half the hash table! no longer constant time...

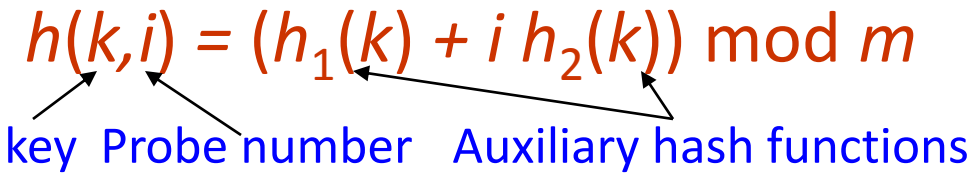
0	49
1	58
2	9
3	
4	
5	
6	
7	
8	18
9	89

Quadratic Probing

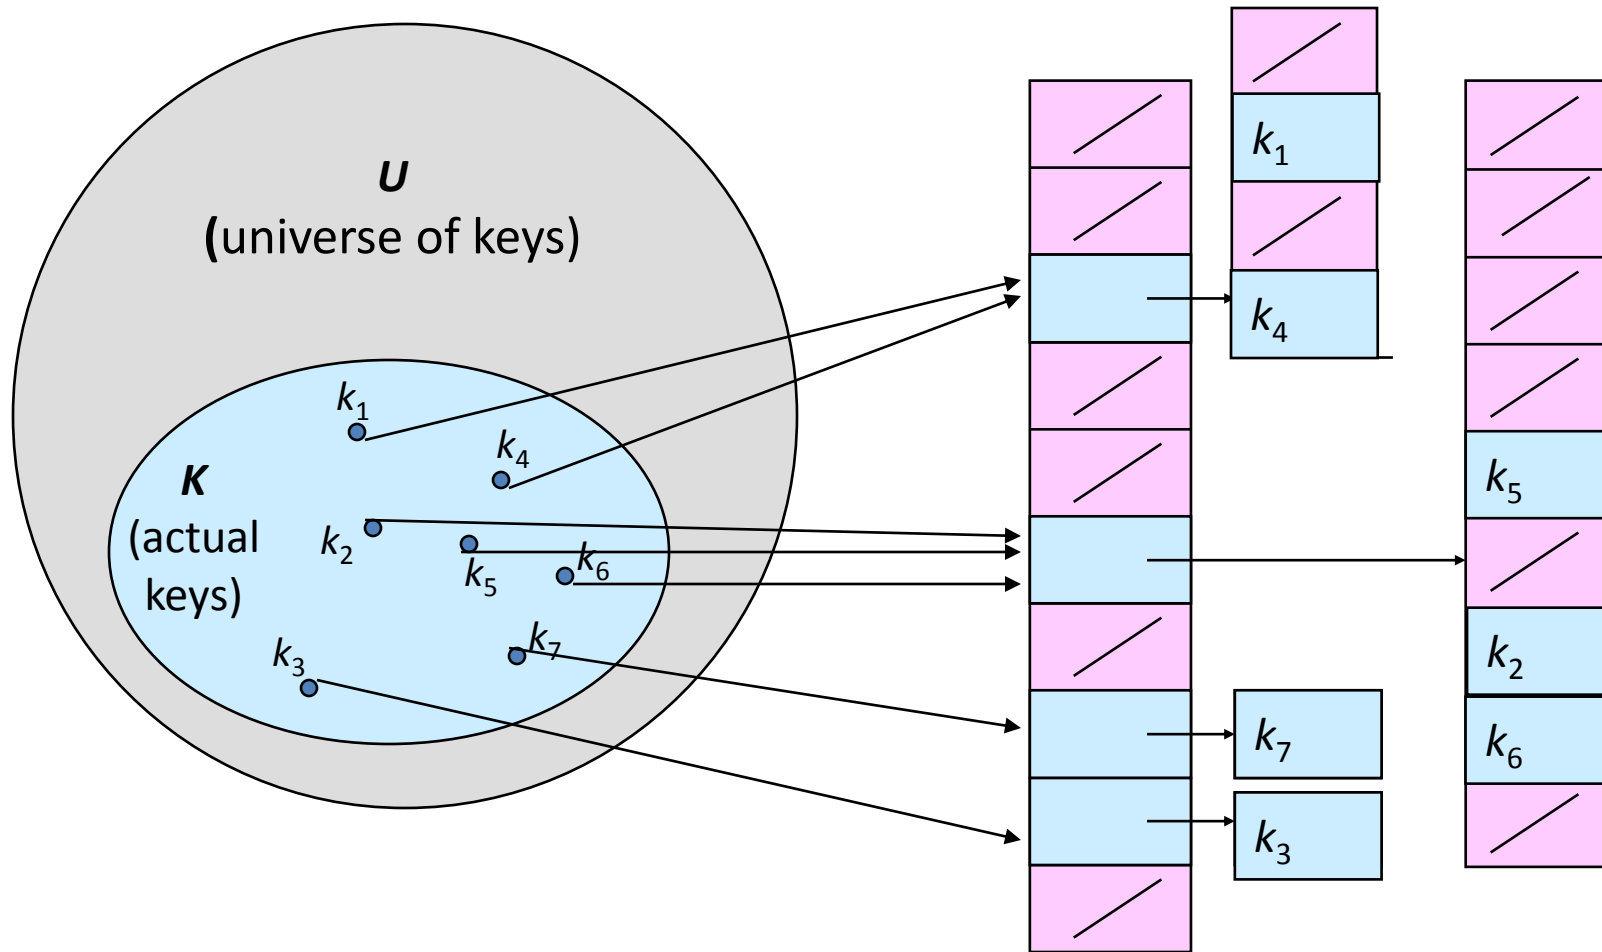
- $h(k,i) = (h'(k) + c_1i + c_2i^2) \bmod m$ $c_1 \neq c_2$
- Can suffer from *secondary clustering*
- **Example:** resolving collisions on slot i by putting the colliding element into slot $i+1, i+4, i+9, i+16, \dots$
 - add 89, 18, 49, 58, 9
 - 49 collides (89 is already there), so we search ahead by +1 to empty slot 0
 - 58 collides (18 is already there), so we search ahead by +1 to occupied slot 9, then +4 to empty slot 2
 - 9 collides (89 is already there), so we search ahead by +1 to occupied slot 0, then +4 to empty slot 3
 - clustering is reduced
 - what is the lookup algorithm?

0	49
1	
2	58
3	9
4	
5	
6	
7	
8	18
9	89

Double Hashing


- $h(k,i) = (h_1(k) + i h_2(k)) \bmod m$

key Probe number Auxiliary hash functions
- Two auxiliary hash functions.
 - h_1 gives the initial probe. h_2 gives the remaining probes.
- Must have $h_2(k)$ relatively prime to m , so that the probe sequence is a full permutation of $\langle 0, 1, \dots, m-1 \rangle$.
 - Choose m to be a power of 2 and have $h_2(k)$ always return an odd number. Or,
 - Let m be prime, and have $1 < h_2(k) < m$.
- $\Theta(m^2)$ different probe sequences.
 - One for each possible combination of $h_1(k)$ and $h_2(k)$.
 - Close to the ideal uniform hashing.

Perfect Hashing



Perfect Hashing

- If you know the n keys in advance, makes a hash table with $O(n)$ size, and worst-case $O(1)$ lookup time.
- Just use two levels of hashing:
A table of size n , then tables of size n_j^2 .
- Dynamic versions have been created, but are usually less practical than other hash methods.
- Key idea: exploit both ends of space/#collisions tradeoff.



Analysis of hash tables

- Main operation: lookup of item in table
- What is worst-case cost of finding an item?
- Is the worst-case cost different for chaining, and the various open addressing schemes?
- Worst-case analysis doesn't make sense for hash tables, look at average case cost
- Cost highly depend on the **load factor** (no. of elements / array size)
- Which is better – hashing or tree based representation?